



Building reliable and high
performance messaging
applications with POJOs
and Apache ActiveMQ

James Strachan

Personal Introduction

- James Strachan
 - Apache Member & Committer on many open source projects at
 - Apache, Codehaus, OpenSymphony, SourceForge, Tigris
 - Co-founder of a few Open Source Projects
 - Apache ActiveMQ, Apache ServiceMix, Groovy, Apache Geronimo, Apache Ode, dom4j, Jaxen etc.
 - Chief Architect, LogicBlaze
 - An Enterprise Open Source company
 - Provide training, mentoring, support for open source SOA

- Messaging 101
- Overview of Apache ActiveMQ
- Advanced features of Apache ActiveMQ
- Developing JMS applications with POJOs
- Conclusions

- **Messaging is...**
 - Loosely coupled exchange of messages between producers and consumers
 - So producers and consumers know nothing of each other
 - they only know about destinations (queues and topics)
 - the ideal approach to building high performance distributed systems
 - Web services done right is really just a form of messaging using pointy brackets
 - Can be persistent or non-persistent
 - May add timeouts & priorities to messages

Topics (publish & subscribe)

- One message goes to 0-to-many consumers based on the current subscribers
 - Think like mailing lists or discussion forums
 - The producer is decoupled from the consumers; it doesn't need to know who all the consumers are
- Ideal for publishing business events
 - Distributed observer pattern
 - Allows one part of your system to notify anyone else who may be interested in an event

Queues (load balancing)

- Messages are load balanced across many consumers
 - Each message goes to exactly one consumer
 - Consumers compete for messages
- Its easy to browse and monitor queues
 - Monitor the performance of your application, find the hotspots, alert if queues are too full
- If a consumer crashes during the processing of a message it is automatically redelivered to another consumer
- So queues implement reliable load balancing with optional persistence

Queues - ideal for grid style

- Ideal for building reliable grid applications
 - Fire requests at a cluster of servers
 - Let the JMS provider provide load balancing, redelivery and optional persistence
 - If the queue starts filling up, just boot up new servers
 - Easily monitor the performance and status of the system
 - View the throughput rates on each queue and queue size
 - Deals easily with parts of your system being taken down for maintenance

- What is Apache ActiveMQ?

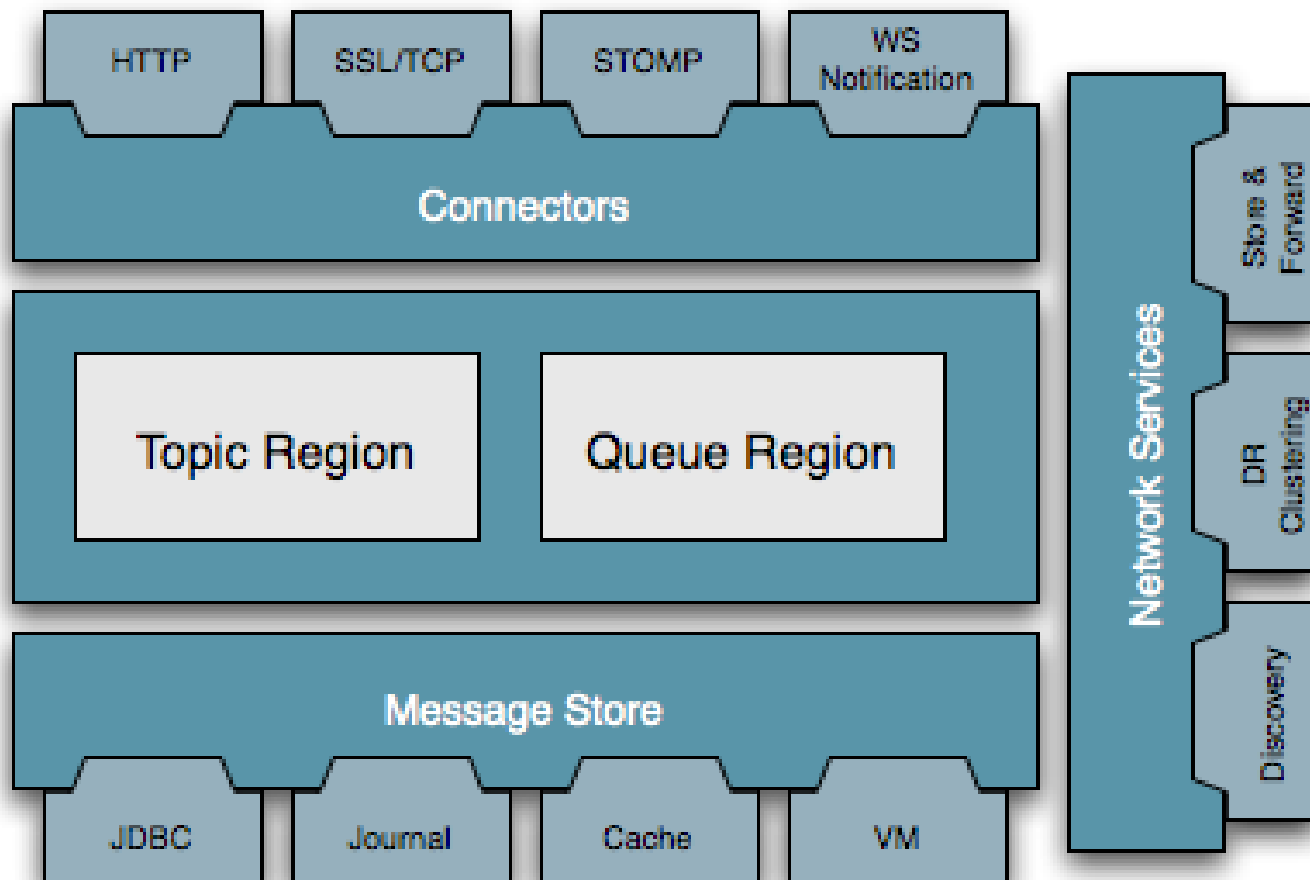
Apache ActiveMQ

- Message fabric:
 - Message broker and clients for many languages and platforms
- Standards-based:
 - JMS 1.1, J2EE 1.4, JCA 1.5 and XA
- Integrated to:
 - JSE, Geronimo, Spring, Tomcat, JBoss and any J2EE 1.4 container (e.g., WebLogic/WebSphere)
- Transport Methods Supported:
 - TCP, SSL, HTTP, OpenWire, Stomp, UDP, multi-cast
- Capabilities include:
 - Queues, Topics, Durable and Non-durable Messaging with Plug-able Security and Persistence



Why Choose? **ActiveMQ**

- **Fast**
 - Fastest open source JMS provider by some margin and close, or better than, the proprietary alternatives
- **Highly scalable**
 - Clustering, peer-to-peer and federated networks support
 - Distributed Destinations
- **Lots of open source clients**
 - Clients for Java, C, C++, C#, ruby, python, perl, php, pike
- **Easy to Use**
 - Minimal Configuration (Dynamic Destination Creation)
- **Open Source**
 - Apache project & Geronimo JMS provider
 - backed by full commercial support from LogicBlaze
- **Lots of advanced features**
 - Covered in the next few slides :-)



- Pluggable features and strategies
 - Transports
 - TCP, SSL, HTTP, UDP, multicast
 - BIO, NIO, AIO
 - Pluggable Wire formats
 - OpenWire, XML, Stomp
 - Security
 - JAAS/JAAS plugin

- JDBC provider
 - Defaults to embedded Apache Derby
 - Can use any RDMBS:
 - MySQL, PostGresql, Oracle, Sybase, DB2, Informix etc
- High performance journal
 - For short term persistence
 - Checkpoints to long term storage (JDBC) at regular intervals
- Kaha provider
 - Complete persistence option without using JDBC which is 2x–4x faster
 - Similar to Berkeley DB in some ways
 - Uses the journal for high performance persistence
 - Updates a separate index file asynchronously

Master / Slave (HA)

- A logical broker is created from a pair of physical brokers
 - a Master and Slave
- All messages and acknowledgements replicated to both physical brokers
 - provides redundancy on hardware failure
 - Auto-failover from master to slave
- Avoids loss of messages even if you have a catastrophic hardware failure or loose a data centre
 - Works on commodity hardware

Networks (store & forward)

- **Federated network of brokers**
 - Brokers have 'network connectors' between each other
 - Store forward or demand based forwarding
 - Can use fixed network routing tables or discovery
 - Clients can either piggy back the discovery mechanism or use fixed host lists etc
- **Persistence model**
 - Each broker owns a message
 - Can store and forward to other brokers
 - So the message 'moves' from broker to broker

Configuration

- Use Java code to configure JMS client/broker
- URI configuration mechanism
 - **vm://localhost?broker.persistent=false**
 - **tcp://localhost:61616?jms.useAsyncSend=true**
- JNDI provider (via properties file)
 - Allows queues/topics to be configured in JNDI
 - Dynamic queue/topic lookup.
 - E.g. client looks up in JNDI to 'dynamicQueues/FOO.BAR' will auto-create the queue dynamically in JNDI
- Spring based XML configuration file & factory beans
 - Easily integrate your own POJOs/extensions into client or broker
- Extended Spring XML configuration file using XBean to make XML more concise & easily validated

- Advanced features of Apache ActiveMQ

Advanced features of ActiveMQ

➤ Wildcards

- Subscribe to different hierarchies in your destination namespaces
 - Products.Books.Computing.EIP
 - Products.Books.Computing.*
 - Products.>

➤ Selectors

- Provide content based filtering on messages using SQL 92 syntax
 - Customer = 'gold' and product in (1, 2, 3) and JMSPriority > 5
- Supports XPath on the message body for XML messages

Advanced features of ActiveMQ

- Durable Topics kinda suck
 - Only one thread allowed to subscribe to a logical subscription
- Virtual Destinations to the rescue!
 - Publish to a topic
 - VirtualTopics.Foo
 - Subscribe to a queue which logically represents a durable topic subscription
 - Consumer.MyConsumerName.VirtualTopics.Foo
 - Now can have many consumers load balancing
 - Can browse the subscription like any other queue

Advanced features of ActiveMQ

- Easy unit testing of JMS
 - Really fast startup
 - No separate process required
 - No persistence so no need to clear out queues before/after tests

```
ConnectionFactory factory =  
    new ActiveMQConnectionFactory("vm://localhost?broker.persistent=false");
```


Advanced features of ActiveMQ

- Embedded brokers
 - Run the message broker inside the JVM of a producer or consumer
 - Can reduce a network hop
- Supports peer based and federated networks
 - Peer transport
 - JMS clients auto discover each other and use embedded brokers

Advanced features of ActiveMQ

- Automatic reconnection
 - If a broker dies or a network error prevents connection, the client can
 - automatically reconnect to another available broker
 - redeliver any in-flight messages
 - Customizable back-off algorithm etc
 - Pluggable discovery mechanism or
 - Fixed list of hosts/URLs to use
 - Clients can use failover: URL to connect to one of a number of brokers

failover:tcp://host1:port,tcp://host2:port,http://host3:port

Advanced features of ActiveMQ

- Queue ordering
 - Order is maintained for messages on a queue for a single consumer
 - With any JMS provider if you have multiple consumers on a queue you lose ordering
- One solution is Exclusive Queues
 - Basically one consumer owns the queue until it dies, then another consumer takes over
 - Useful if you want to deploy a homogeneous cluster of consumers where queue ordering is important and you want to auto-failover if a node dies
 - Though there is no parallelisation with this approach
 - Just 1 thread is processing all of the messages

Advanced features of ActiveMQ

- Message Groups (sticky load balancing & grid partitioning)
 - Maintain order for messages within a Message Group while providing load balancing across a cluster of consumers for different Message Groups
 - A Message Group is defined by setting the JMSXGroupID String header
 - E.g. set JMSXGroupID to the user ID or product code
 - All messages for that user/product will maintain their order and be processed usually by the same consumer
 - Different users/products will be processed in parallel on different consumers
 - Implements sticky load balancing for JMS
 - JMSXGroupID is a little like jsessionid in the web world
 - Efficiently partition your messages across many consumers
 - Reliable and supports failover
 - each value of JMSXGroupID is owned by one consumer
 - If the consumer dies, another one takes over and messages are redelivered
 - Preserves ordering but allows increased parallelisation and load balancing
 - Reduces concurrency and increases cache efficiency
 - Only 1 thread in your network will be processing a single JMSXGroupID value at once

Multi-client Connectivity

- **REST**
 - Use regular HTTP operations
- **Ajax Messaging**
 - very efficient Comet server using ActiveMQ and Jetty
- **Stomp**
 - Current: Ruby, Python, Perl, PHP, .Net/C#, C, C++
 - Note that the Ruby, Python, Perl and PHP clients are pure – no C library required
- **OpenWire**
 - works for Java, C, C++ and C# today (all pure language clients)
- **ServiceMix**
 - Multi-protocol bridge (e.g. MQSeries/RV to ActiveMQ bridge etc.)
 - Supported clients/transport: email, jabber, FTP, WebDAV, Samba, HTTP, foreign JMS
 - WS-Notification & WS-ReliableMessaging (via ServiceMix)
 - any JBI/JCA binding (e.g. CORBA/CICS/SAP)

- Designed for high performance
 - Binary & small and fast as possible
 - Can go beyond Stomp/JMS in features
 - Backwards compatible with any 4.x client
 - So can upgrade broker and not worry about clients
- Code generate marshalling code & test suite
 - Use Java Command POJOs with annotations
 - Guarantees all clients work at the marshalling level
- ActiveMQ 4.0 has full OpenWire support for
 - Java, C, C++, C#
 - Both C++ and C# have their own language versions of the JMS API (CMS and NMS respectively) which are easy to implement on any transport (e.g. Stomp and OpenWire)

Stomp overview

- Designed so anyone can write a client really easily in any language
- Supports core JMS/MOM features
 - Connection, security
 - Sending, Receiving, Subscribe & Unsubscribe
 - Concurrent subscriptions, requests and response
 - Message acknowledgements (or auto-ack)
 - Transactions
 - JMS headers (correlation, replyto, message-id, message group etc)
 - Different QoS
 - queues, topics, temporary queue, topics
 - durable/non-durable)
- Optimized for ease of client (so makes a few performance trade offs)
 - E.g. text based, like HTTP
 - A little bigger than need be
 - A little slower than could be to parse
 - Default is text based rather than content-size/binary based

Stomp: connecting

These examples show using telnet as a Stomp client...

Sends...

```
CONNECT
login:<username>
passcode:<passcode>
```

^@

Receives

```
CONNECTED
session:<session-id>
```

^@

Stomp: sending messages

Send a message to a queue or topic

```
SEND
```

```
destination:/topic/foo.bar
```

```
hello world!
```

```
^@
```

Stomp: sending with headers

Send using correlation, request/response, headers

```
SEND
```

```
destination:/queue/orders.books
```

```
reply-to:/temporaryQueue/James
```

```
correlation-id:4324234
```

```
amazonSecurityToken:abc234
```

```
amazonCustomerRating:gold
```

```
<order id="123" customer="jstrachan">
```

```
  <book isin="1234"/>
```

```
  <book isin="456"/>
```

```
</order>
```

```
^@
```

Stomp: subscribing

Subscribe first

```
SUBSCRIBE  
destination:/queue/orders.books
```

```
^@
```

Then later you'll receive inbound messages

```
MESSAGE  
destination:/queue/orders.books  
Message-id:abc123  
reply-to:/temporaryQueue/james134
```

This is the message

```
^@
```

- Developing JMS applications with POJOs

Using JMS efficiently

- Pool expensive resources like Connections, Sessions, MessageProducers, MessageConsumers
 - They are relatively expensive to create/destroy as each create/destroy requires a blocking request/response with the broker
 - These objects are designed to be super efficient after they are created, so create them on startup and pool
- Be very careful with Spring's JmsTemplate
 - Each operation creates/destroys a Connection, Session, MessageProducer/MessageConsumer
 - Normally very inefficient unless you use a pooling wrapper over your JMS provider
 - E.g. PoolingConnectionFactory in ActiveMQ
 - Only good for sending

Using JMS efficiently

- Synchronous v Asynchronous trade-off
 - Asynchronous is much faster; though sometimes for reliability you must block until a message is on disk
- JMS transactions boost performance when performing multiple operations
 - E.g. consume a message and send 1..N messages in a transaction
 - Only pay the *sync cost* once, on the commit/rollback rather than on each send() or acknowledge()
 - With XA you can use batching to reduce the *sync cost*
- ActiveMQ Tuning Guide
 - <http://devzone.logicblaze.com/site/how-to-tune-activemq.html>

Using JMS efficiently

- Consuming messages in servers efficiently is harder than you might think
 - Each JMS Session uses 1 thread to process messages for its MessageConsumers
 - This is usually fine for GUIs
 - In servers you typically want a large thread pool and pool of Connections/Sessions/MessageConsumers to process inbound messages efficiently
 - You often want a pool of MessageListener objects too
- Message Driven Beans are one solution
 - They are EJBs and use JCA to do the JMS resource & thread pooling together with transactions & exception handling

- Write POJOs for your business logic
 - Keep your POJOs free of any middleware code
- Inject the middleware at deployment time
 - E.g. Spring Remoting, SCA, JAX-WS
 - Clean separation of concerns between middleware and business logic
- Middleware is increasingly becoming invisible!
 - Allows you to change the middleware easily as your requirements and topology change

- <http://lingo.codehaus.org/>
- POJO Remoting for JMS using Spring
 - Hides the JMS code from your application
 - uses dynamic proxies/cglib to make client proxies
- Supports various *message exchange patterns*
 - Synchronous request/response
 - Asynchronous one ways & subscriptions
 - Asynchronous request/response
- Works well with JCA and Jencks

```
public interface ExampleService {  
  
    Foo regularRPC(String name);  
    void anotherRPC() throws Exception;  
  
    @OneWay  
    void someOneWayMethod(String name, int  
        age);  
  
    @OneWay  
    void watchPrices(String stock, MyCallback  
        listener);  
}
```


Lingo - client side

```
<!-- client side proxy-->
<bean id="client"
      class="org.logicblaze.lingo.jms.JmsProxyFactoryBean">
  <property name="serviceInterface" value="com.acme.ExampleService"/>
  <property name="connectionFactory" ref="jmsFactory"/>
  <property name="destination" ref="exampleDestination"/>
</bean>

<!-- JMS ConnectionFactory to use -->
<bean id="jmsFactory"
      class="org.apache.activemq.ActiveMQConnectionFactory">
  <property name="brokerURL" value="tcp://localhost:61616"/>
</bean>

<!-- JMS Destination to use -->
<bean id="exampleDestination"
      class="org.apache.activemq.command.ActiveMQQueue">
  <constructor-arg index="0"
    value="test.org.logicblaze.lingo.example"/>
</bean>
```

Lingo - server side

```
<!-- the server side -->
<bean id="server"
      class="org.logicblaze.lingo.jms.JmsServiceExporte
r">
  <property name="service" ref="serverImpl"/>
  <property name="serviceInterface"
    value="com.acme.ExampleService"/>
  <property name="connectionFactory"
    ref="jmsFactory"/>
  <property name="destination"
    ref="exampleDestination"/>
</bean>

<!-- implementation of the service -->
<bean id="serverImpl"
      class="com.acme.ExampleServiceImpl"
      singleton="true"/>
```

➤ **Conclusions**

Conclusions

- Apache ActiveMQ is a great way to build a high performance, reliable and scalable distributed system
 - Load balancing
 - Data partitioning
 - High Availability
 - Loose coupling
- POJO frameworks like Lingo and Jencks can help hide the JMS plumbing code and implement efficient pooling and concurrent processing of messages

➤ Useful links

➤ Apache ActiveMQ

➤ <http://incubator.apache.org/activemq/>

➤ Apache ServiceMix (JBI based ESB)

➤ <http://incubator.apache.org/servicemix/>

➤ Jencks

➤ <http://jencks.codehaus.org/>

➤ Lingo

➤ <http://lingo.codehaus.org/>

➤ Support and Services

➤ [Http://logicblaze.com](http://logicblaze.com)

➤ Questions and Answers

The logo for Apache ActiveMQ, featuring the word "ActiveMQ" in blue, with a stylized blue swoosh above the text.The logo for Apache ServiceMix, featuring the word "ServiceMix" in blue, with a stylized blue circular swoosh above the text.

- <http://jencks.codehaus.org/>
- Implements Message Driven POJOs
 - Local and XA Transaction handling
 - Pools connections, sessions, threads
 - Exception handling & retry etc
 - A full JCA container which reuses the JCA Resource Adapters
- Supports any API – JMS, JDBC, JCA CCI, JBI, JAX-RPC etc
- Supports outbound pooling
 - E.g. pooled JMS sending, JDBC
 - Works with local or XA transactions too
- Uses the JCA Resource Adapters for resources for optimal pooling, resource usage and performance
 - E.g. the ActiveMQ Resource Adapter offers sophisticated concurrency control, resource usage and batching of requests to boost performance when using XA


```
<!-- JCA Container -->
<bean id="jencks" class="org.jencks.JCAContainer">
  <property name="bootstrapContext">
    <bean
      class="org.jencks.factory.BootstrapContextFactoryBean">
        <property name="threadPoolSize" value="25"/>
      </bean>
    </property>

    <property name="resourceAdapter">
      <bean id="activeMQResourceAdapter"
        class="org.apache.activemq.ra.ActiveMQResourceAdapter">
        <property name="serverUrl"
          value="tcp://localhost:61616"/>
        </bean>
      </property>
    </bean>
  </bean>
```

```
<!-- Message Driven POJO -->
<bean id="inboundConnectorA" class="org.jencks.JCAConnector">
  <property name="jcaContainer" ref="jencks" />

  <!-- subscription details -->
  <property name="activationSpec">
    <bean class="org.apache.activemq.ra.ActiveMQActivationSpec">
      <property name="destination" value="Foo.Bar"/>
      <property name="destinationType" value="javax.jms.Topic"/>
    </bean>
  </property>

  <property name="ref" value="echoBean"/>
</bean>

<!-- The actual POJO for processing messages which implements
MessageListener -->
<bean id="echoBean" class="com.acme.MyMessageListener"
  singleton="true"/>
```

- ActivationSpec allows you to configure
 - Destination & optional selector and durable topic subscriptions
- ActiveMQActivationSpec extensions
 - Concurrency per subscription
 - how many messages can be processed in parallel
 - Redelivery policy
 - If a message fails use exponential backoff?
 - How many redeliveries?
 - Depends on what the listener does on how you deal with it
 - is a particular operation prone to deadlocks?
 - Pre-fetch options
 - How aggressively should the Resource Adapter preload asynchronously messages to be processed
 - RAM versus performance tradeoffs